

Yao Yao Patientfi Data Challenge

July 31, 2022

1 Yao Yao

Patientfi Data Challenge

- 1.0.1 You will find a file labeled data.csv in the link that contains loan data including the credit attributes and loan status, etc. Your goal is to predict the likelihood of a loan to “charge off” (column “loan_status”).
- 1.0.2 The credit default prediction is central to managing risk in a consumer lending business, as it allows lenders to optimize lending decisions. You’ll apply your machine learning skills to predict credit default by the dataset we provided to you. You’re free to explore any techniques to create the most powerful model, from creating features to using the data in a more organic way within a model.
- 1.0.3 Please build your model with a detailed description/explanation, as it will help us understand and evaluate how you approach the problem.

```
[1]: import sys
try:
    sys.getwindowsversion()
except AttributeError:
    isWindows = False
else:
    isWindows = True
if isWindows:
    import win32api, win32process, win32con
    pid = win32api.GetCurrentProcessId()
    handle = win32api.OpenProcess(win32con.PROCESS_ALL_ACCESS, True, pid)
    win32process.SetPriorityClass(handle, win32process.HIGH_PRIORITY_CLASS)
import warnings
warnings.filterwarnings('ignore')

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
from sklearn.ensemble import RandomForestClassifier, RandomForestRegressor
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix, roc_auc_score, roc_curve
from sklearn.tree import DecisionTreeClassifier
from sklearn.preprocessing import MinMaxScaler
from sklearn import metrics
from datetime import datetime
import time
from sklearn.model_selection import KFold, StratifiedKFold
import itertools
import os
from PIL import Image
os.environ["PATH"] += os.pathsep + 'C:/Program Files (x86)/Graphviz2.38/bin/'

%matplotlib inline

```

2 Import data and create data summary

```
[2]: df=pd.read_csv("data.csv", engine = 'c')
```

```
[3]: df.describe().T
```

```
[3]:
```

	count	mean	std	\
id	199999.0	6.229635e+07	3.941126e+06	
member_id	0.0	NaN	NaN	
loan_amnt	199999.0	1.527812e+04	8.651139e+03	
funded_amnt	199999.0	1.527812e+04	8.651139e+03	
funded_amnt_inv	199999.0	1.526947e+04	8.646325e+03	
...	
hardship_payoff_balance_amount	1409.0	1.001451e+04	6.313702e+03	
hardship_last_payment_amount	1409.0	1.860614e+02	1.812045e+02	
settlement_amount	5840.0	5.015406e+03	3.557532e+03	
settlement_percentage	5840.0	4.710027e+01	5.753855e+00	
settlement_term	5840.0	1.348545e+01	7.537550e+00	

	min	25%	50%	\
id	56705.00	59411733.00	62217538.00	
member_id	NaN	NaN	NaN	
loan_amnt	1000.00	8500.00	14000.00	
funded_amnt	1000.00	8500.00	14000.00	
funded_amnt_inv	900.00	8475.00	14000.00	
...	
hardship_payoff_balance_amount	55.73	4708.77	8879.16	
hardship_last_payment_amount	0.02	44.95	134.56	
settlement_amount	130.00	2192.49	4344.99	
settlement_percentage	20.00	45.00	45.00	

settlement_term	0.00	8.00	13.50
-----------------	------	------	-------

	75%	max
id	65644568.00	68617057.00
member_id	NaN	NaN
loan_amnt	20000.00	35000.00
funded_amnt	20000.00	35000.00
funded_amnt_inv	20000.00	35000.00
...
hardship_payoff_balance_amount	14383.87	29401.04
hardship_last_payment_amount	275.29	927.79
settlement_amount	7000.00	26242.50
settlement_percentage	50.00	97.66
settlement_term	18.00	65.00

[115 rows x 8 columns]

```
[4]: df.isnull().T.any().T.sum()
```

```
[4]: 199999
```

3 Given that there are a lot of values missing, remove columns that have over 80% missing for the sake of the exercise and time. 'Default' for loan status is removed because it only has one result

```
[5]: df = df[df.columns[df.isnull().mean() < 0.8]]
df = df[df.columns[df.isna().mean() < 0.8]]
df = df[df['loan_status'] != 'Default'].reset_index(drop=True)
df
```

```
[5]:
```

	id	loan_amnt	funded_amnt	funded_amnt_inv	term \
0	68407277	3600	3600	3600.0	36 months
1	68355089	24700	24700	24700.0	36 months
2	68341763	20000	20000	20000.0	60 months
3	66310712	35000	35000	35000.0	60 months
4	68476807	10400	10400	10400.0	60 months
...
199993	56059770	4000	4000	4000.0	36 months
199994	56080425	12000	12000	12000.0	36 months
199995	55909672	21000	21000	21000.0	36 months
199996	54414556	27500	27500	27500.0	60 months
199997	56109383	7000	7000	7000.0	36 months

	int_rate	installment	grade	sub_grade	emp_title \
0	13.99	123.03	C	C4	leadman

1	11.99	820.28	C	C1	Engineer
2	10.78	432.66	B	B4	truck driver
3	14.85	829.90	C	C5	Information Systems Officer
4	22.45	289.91	F	F1	Contract Specialist
...
199993	12.29	133.42	C	C1	Teacher
199994	12.69	402.54	C	C2	Truck driver
199995	12.29	700.42	C	C1	Attorney
199996	14.65	649.19	C	C5	hvac tech
199997	10.99	229.14	B	B4	Nursing Office Manager

	...	percent_bc_gt_75	pub_rec_bankruptcies	tax_liens	tot_hi_cred_lim	\
0	...	0.0	0	0	178050	
1	...	7.7	0	0	314017	
2	...	50.0	0	0	218418	
3	...	0.0	0	0	381215	
4	...	60.0	0	0	439570	
...	
199993	...	83.3	0	0	194282	
199994	...	50.0	0	0	32176	
199995	...	100.0	0	0	181446	
199996	...	0.0	0	0	53653	
199997	...	100.0	0	0	134023	

	total_bal_ex_mort	total_bc_limit	total_il_high_credit_limit	\
0	7746	2400	13734	
1	39475	79300	24667	
2	18696	6200	14877	
3	52226	62500	18000	
4	95768	20300	88097	
...	
199993	29295	24400	7000	
199994	27413	13800	5784	
199995	173683	15000	157346	
199996	30750	18500	35153	
199997	38802	4200	22492	

	hardship_flag	disbursement_method	debt_settlement_flag
0	N	Cash	N
1	N	Cash	N
2	N	Cash	N
3	N	Cash	N
4	N	Cash	N
...
199993	N	Cash	N
199994	N	Cash	N
199995	N	Cash	N

199996	N	Cash	N
199997	N	Cash	N

[199998 rows x 94 columns]

```
[6]: df.isnull().T.any().T.sum()
```

```
[6]: 176141
```

```
[7]: #df.to_csv('df.csv',index=False)
```

4 Convert certain strings into continuous variables

```
[8]: df.describe().T
```

```
[8]:
```

	count	mean	std	min	\
id	199998.0	6.229633e+07	3.941125e+06	56705.00	
loan_amnt	199998.0	1.527802e+04	8.651048e+03	1000.00	
funded_amnt	199998.0	1.527802e+04	8.651048e+03	1000.00	
funded_amnt_inv	199998.0	1.526937e+04	8.646234e+03	900.00	
int_rate	199998.0	1.236171e+01	4.242075e+00	5.32	
...	
tax_liens	199998.0	6.433564e-02	4.656492e-01	0.00	
tot_hi_cred_lim	199998.0	1.770256e+05	1.779942e+05	2500.00	
total_bal_ex_mort	199998.0	5.281753e+04	4.949148e+04	0.00	
total_bc_limit	199998.0	2.258573e+04	2.234689e+04	0.00	
total_il_high_credit_limit	199998.0	4.413560e+04	4.447859e+04	0.00	
		25%	50%	75%	max
id	59411731.00	62217530.00	65644565.25	68617057.00	
loan_amnt	8500.00	14000.00	20000.00	35000.00	
funded_amnt	8500.00	14000.00	20000.00	35000.00	
funded_amnt_inv	8475.00	14000.00	20000.00	35000.00	
int_rate	9.17	12.29	14.65	28.99	
...	
tax_liens	0.00	0.00	0.00	85.00	
tot_hi_cred_lim	52205.00	115187.50	255110.50	999999.00	
total_bal_ex_mort	22736.00	40069.00	66654.50	2652799.00	
total_bc_limit	8000.00	15700.00	29500.00	834300.00	
total_il_high_credit_limit	15850.25	33575.50	59213.00	2101913.00	

[74 rows x 8 columns]

5 Convert term into continuous variable

```
[9]: df['term'] = df['term'].str[:3].astype('int')
df['term']
```

```
[9]: 0      36
     1      36
     2     60
     3     60
     4     60
     ..
199993  36
199994  36
199995  36
199996  60
199997  36
```

Name: term, Length: 199998, dtype: int32

6 Make Grade numeric multiplied by 10, and Sub Grade numeric multiplied by 2 minus 2 to keep it at scale

```
[10]: df['grade'] = df['grade'].to_numpy().astype('<U1').view(np.uint32)-64
df['grade']
```

```
[10]: 0      3
     1      3
     2      2
     3      3
     4      6
     ..
199993  3
199994  3
199995  3
199996  3
199997  2
```

Name: grade, Length: 199998, dtype: uint32

```
[11]: df['sub_grade'] = df['sub_grade'].str[1:].astype('int')*2-2
df['sub_grade']
```

```
[11]: 0      6
     1      0
     2      6
     3      8
     4      0
     ..
```

```

199993    0
199994    2
199995    0
199996    8
199997    6
Name: sub_grade, Length: 199998, dtype: int32

```

```
[12]: df['grade'] = df['grade']*10+df['sub_grade']
df['grade']
```

```

[12]: 0          36
      1          30
      2          26
      3          38
      4          60
      ..
199993    30
199994    32
199995    30
199996    38
199997    26
Name: grade, Length: 199998, dtype: int64

```

7 Make the titles lowercase, trim whitespace, and then make dummy variables for the words separated by space later

```
[13]: df['emp_title'] = df['emp_title'].str.lower().str.strip()
df['emp_title']
```

```

[13]: 0          leadman
      1          engineer
      2      truck driver
      3  information systems officer
      4      contract specialist
      ...
199993          teacher
199994      truck driver
199995          attorney
199996          hvac tech
199997  nursing office manager
Name: emp_title, Length: 199998, dtype: object

```

make the length of the career numeric, simplified

8 Change length of title to numerical

```
[14]: df['emp_length'] = df['emp_length'].str.split(' ', 1, expand=True)[0].str.  
      ↪extract('(\d+)', expand=False).replace('None', np.nan).astype('float')  
      df['emp_length']
```

```
[14]: 0          10.0  
      1          10.0  
      2          10.0  
      3          10.0  
      4           3.0  
      ...  
      199993      10.0  
      199994       1.0  
      199995       8.0  
      199996      10.0  
      199997       7.0  
      Name: emp_length, Length: 199998, dtype: float64
```

9 Zip code can be converted into numeric given there that is a pattern in geo labeling

```
[15]: df['zip_code'] = df['zip_code'].str[:3].replace('None', np.nan).astype('int')  
      df['zip_code']
```

```
[15]: 0          190  
      1          577  
      2          605  
      3           76  
      4          174  
      ...  
      199993      850  
      199994      983  
      199995      113  
      199996      983  
      199997      846  
      Name: zip_code, Length: 199998, dtype: int32
```

10 Convert the dates into datetime format

```
[16]: df['earliest_cr_line'] = pd.to_datetime(df['earliest_cr_line'], format='%b-%y')  
      df['earliest_cr_line']
```

```
[16]: 0          2003-08-01  
      1          1999-12-01  
      2          2000-08-01
```



```

3          2008-09-01
4          1998-06-01
...
199993     1997-06-01
199994     2003-09-01
199995     2001-10-01
199996     2002-06-01
199997     1990-11-01
Name: earliest_cr_line, Length: 199998, dtype: datetime64[ns]

```

```
[17]: df['earliest_cr_line'].min()
```

```
[17]: Timestamp('1969-01-01 00:00:00')
```

```
[18]: df['last_credit_pull_d'] = pd.to_datetime(df['last_credit_pull_d'],
        ↪format='%b-%y')
df['last_credit_pull_d']
```

```

[18]: 0          2019-03-01
1          2019-03-01
2          2019-03-01
3          2019-03-01
4          2018-03-01
...
199993     2019-02-01
199994     2016-02-01
199995     2018-09-01
199996     2016-02-01
199997     2019-03-01
Name: last_credit_pull_d, Length: 199998, dtype: datetime64[ns]

```

```
[19]: df['last_pymnt_d'] = pd.to_datetime(df['last_pymnt_d'], format='%b-%y')
df['last_pymnt_d']
```

```

[19]: 0          2019-01-01
1          2016-06-01
2          2017-06-01
3          2019-02-01
4          2016-07-01
...
199993     2018-06-01
199994     2016-01-01
199995     2018-08-01
199996     2016-03-01
199997     2018-08-01
Name: last_pymnt_d, Length: 199998, dtype: datetime64[ns]

```

11 Earliest credit line, credit pull, and last payment dates are less arbitrary when relative to each other, so I made new columns showing the difference in days among those dates

```
[20]: df['diff credit pull credit line'] = (df['last_credit_pull_d'] -  
      ↪ df['earliest_cr_line']).dt.days  
df['diff credit pull credit line']
```

```
[20]: 0          5691.0  
      1          7030.0  
      2          6786.0  
      3          3833.0  
      4          7213.0  
      ...  
      199993      7915.0  
      199994      4536.0  
      199995      6179.0  
      199996      4993.0  
      199997     10347.0  
      Name: diff credit pull credit line, Length: 199998, dtype: float64
```

```
[21]: df['diff payment credit line'] = (df['last_pymnt_d'] - df['earliest_cr_line']).  
      ↪ dt.days  
df['diff payment credit line']
```

```
[21]: 0          5632.0  
      1          6027.0  
      2          6148.0  
      3          3805.0  
      4          6605.0  
      ...  
      199993      7670.0  
      199994      4505.0  
      199995      6148.0  
      199996      5022.0  
      199997     10135.0  
      Name: diff payment credit line, Length: 199998, dtype: float64
```

```
[22]: df['diff credit pull payment'] = (df['last_credit_pull_d'] -  
      ↪ df['last_pymnt_d']).dt.days  
df['diff credit pull payment']
```

```
[22]: 0          59.0  
      1         1003.0  
      2          638.0  
      3          28.0  
      4          608.0
```

```

...
199993    245.0
199994     31.0
199995     31.0
199996    -29.0
199997    212.0
Name: diff credit pull payment, Length: 199998, dtype: float64

```

12 converted datetime to year month on the same scale

```

[23]: df['earliest_cr_line'] = pd.DatetimeIndex(df['earliest_cr_line']).year * 100 +
      ↪(pd.DatetimeIndex(df['earliest_cr_line']).month - 1)* 100/12
      df['earliest_cr_line']

```

```

[23]: 0          200358.333333
      1          199991.666667
      2          200058.333333
      3          200866.666667
      4          199841.666667
      ...
      199993    199741.666667
      199994    200366.666667
      199995    200175.000000
      199996    200241.666667
      199997    199083.333333
Name: earliest_cr_line, Length: 199998, dtype: float64

```

13 Converted month and year into one continuous variable

```

[24]: df['last_credit_pull_d'] = pd.DatetimeIndex(df['last_credit_pull_d']).year *
      ↪100 + (pd.DatetimeIndex(df['last_credit_pull_d']).month - 1)* 100/12
      df['last_credit_pull_d']

```

```

[24]: 0          201916.666667
      1          201916.666667
      2          201916.666667
      3          201916.666667
      4          201816.666667
      ...
      199993    201908.333333
      199994    201608.333333
      199995    201866.666667
      199996    201608.333333
      199997    201916.666667
Name: last_credit_pull_d, Length: 199998, dtype: float64

```

```
[25]: df['last_pymnt_d'] = pd.DatetimeIndex(df['last_pymnt_d']).year * 100 + (pd.
      ↪DatetimeIndex(df['last_pymnt_d']).month - 1)* 100/12
      df['last_pymnt_d']
```

```
[25]: 0          201900.000000
      1          201641.666667
      2          201741.666667
      3          201908.333333
      4          201650.000000
      ...
      199993      201841.666667
      199994      201600.000000
      199995      201858.333333
      199996      201616.666667
      199997      201858.333333
      Name: last_pymnt_d, Length: 199998, dtype: float64
```

14 Sub grade was combined into grade. funded amount is redundant to loan amount. issue date, and distribution method have no variance, so they are dropped

```
[26]: df = df.drop(columns = [
      ↪['issue_d', 'sub_grade', 'disbursement_method', 'funded_amnt']).fillna(-1)
      df
```

```
[26]:
```

	id	loan_amnt	funded_amnt_inv	term	int_rate	installment	\
0	68407277	3600	3600.0	36	13.99	123.03	
1	68355089	24700	24700.0	36	11.99	820.28	
2	68341763	20000	20000.0	60	10.78	432.66	
3	66310712	35000	35000.0	60	14.85	829.90	
4	68476807	10400	10400.0	60	22.45	289.91	
...	
199993	56059770	4000	4000.0	36	12.29	133.42	
199994	56080425	12000	12000.0	36	12.69	402.54	
199995	55909672	21000	21000.0	36	12.29	700.42	
199996	54414556	27500	27500.0	60	14.65	649.19	
199997	56109383	7000	7000.0	36	10.99	229.14	

	grade	emp_title	emp_length	home_ownership	...	\
0	36	leadman	10.0	MORTGAGE	...	
1	30	engineer	10.0	MORTGAGE	...	
2	26	truck driver	10.0	MORTGAGE	...	
3	38	information systems officer	10.0	MORTGAGE	...	
4	60	contract specialist	3.0	MORTGAGE	...	
...	
199993	30	teacher	10.0	MORTGAGE	...	

199994	32	truck driver	1.0	RENT	...
199995	30	attorney	8.0	RENT	...
199996	38	hvac tech	10.0	RENT	...
199997	26	nursing office manager	7.0	MORTGAGE	...

	tax_liens	tot_hi_cred_lim	total_bal_ex_mort	total_bc_limit	\
0	0	178050	7746	2400	
1	0	314017	39475	79300	
2	0	218418	18696	6200	
3	0	381215	52226	62500	
4	0	439570	95768	20300	
...	
199993	0	194282	29295	24400	
199994	0	32176	27413	13800	
199995	0	181446	173683	15000	
199996	0	53653	30750	18500	
199997	0	134023	38802	4200	

	total_il_high_credit_limit	hardship_flag	debt_settlement_flag	\
0	13734	N	N	
1	24667	N	N	
2	14877	N	N	
3	18000	N	N	
4	88097	N	N	
...	
199993	7000	N	N	
199994	5784	N	N	
199995	157346	N	N	
199996	35153	N	N	
199997	22492	N	N	

	diff credit pull	credit line	diff payment	credit line	\
0		5691.0		5632.0	
1		7030.0		6027.0	
2		6786.0		6148.0	
3		3833.0		3805.0	
4		7213.0		6605.0	
...		
199993		7915.0		7670.0	
199994		4536.0		4505.0	
199995		6179.0		6148.0	
199996		4993.0		5022.0	
199997		10347.0		10135.0	

	diff credit pull	payment
0		59.0
1		1003.0

```

2          638.0
3          28.0
4          608.0
...
199993     245.0
199994      31.0
199995      31.0
199996     -29.0
199997     212.0

```

[199998 rows x 93 columns]

```
[27]: df.isna().any().sum()
```

```
[27]: 0
```

15 Summary of variables after clean up

```
[28]: df.describe()
```

```
[28]:
```

	id	loan_amnt	funded_amnt_inv	term \
count	1.999980e+05	199998.000000	199998.000000	199998.000000
mean	6.229633e+07	15278.018530	15269.374289	43.842678
std	3.941125e+06	8651.048016	8646.233566	11.256878
min	5.670500e+04	1000.000000	900.000000	36.000000
25%	5.941173e+07	8500.000000	8475.000000	36.000000
50%	6.221753e+07	14000.000000	14000.000000	36.000000
75%	6.564457e+07	20000.000000	20000.000000	60.000000
max	6.861706e+07	35000.000000	35000.000000	60.000000

	int_rate	installment	grade	emp_length \
count	199998.000000	199998.000000	199998.000000	199998.000000
mean	12.361711	441.381145	31.132161	5.509750
std	4.242075	247.050693	12.686827	4.106834
min	5.320000	14.770000	10.000000	-1.000000
25%	9.170000	261.880000	22.000000	2.000000
50%	12.290000	383.810000	30.000000	6.000000
75%	14.650000	580.730000	38.000000	10.000000
max	28.990000	1445.460000	78.000000	10.000000

	annual_inc	zip_code	...	percent_bc_gt_75 \
count	1.999980e+05	199998.000000	...	199998.000000
mean	7.815054e+04	507.370864	...	46.015581
std	8.051398e+04	309.216121	...	36.389153
min	0.000000e+00	7.000000	...	-1.000000
25%	4.756650e+04	231.000000	...	11.100000

50%	6.500000e+04	468.000000	...	50.000000
75%	9.340000e+04	797.000000	...	75.000000
max	9.000000e+06	999.000000	...	100.000000

	pub_rec_bankruptcies	tax_liens	tot_hi_cred_lim \
count	199998.000000	199998.000000	1.999980e+05
mean	0.132301	0.064336	1.770256e+05
std	0.384264	0.465649	1.779942e+05
min	0.000000	0.000000	2.500000e+03
25%	0.000000	0.000000	5.220500e+04
50%	0.000000	0.000000	1.151875e+05
75%	0.000000	0.000000	2.551105e+05
max	9.000000	85.000000	9.999999e+06

	total_bal_ex_mort	total_bc_limit	total_il_high_credit_limit \
count	1.999980e+05	199998.000000	1.999980e+05
mean	5.281753e+04	22585.728182	4.413560e+04
std	4.949148e+04	22346.885365	4.447859e+04
min	0.000000e+00	0.000000	0.000000e+00
25%	2.273600e+04	8000.000000	1.585025e+04
50%	4.006900e+04	15700.000000	3.357550e+04
75%	6.665450e+04	29500.000000	5.921300e+04
max	2.652799e+06	834300.000000	2.101913e+06

	diff credit pull credit line	diff payment credit line \
count	199998.000000	199998.000000
mean	7116.756833	6838.951250
std	3029.208281	3054.193244
min	-19390.000000	-19390.000000
25%	5264.000000	4960.000000
50%	6664.000000	6390.000000
75%	8644.000000	8401.000000
max	18321.000000	18262.000000

	diff credit pull payment
count	199998.000000
mean	273.524955
std	313.447481
min	-1308.000000
25%	0.000000
50%	181.000000
75%	457.000000
max	1311.000000

[8 rows x 83 columns]

```
[29]: df[df.columns[df.dtypes == "O"]].describe().T
```

```
[29]:
```

	count	unique	top	freq
emp_title	199998	53276	-1	12130
home_ownership	199998	4	MORTGAGE	99617
verification_status	199998	3	Source Verified	84755
loan_status	199998	6	Fully Paid	140991
pymnt_plan	199998	2	n	199976
addr_state	199998	49	CA	27304
initial_list_status	199998	2	w	142654
application_type	199998	2	Individual	199487
hardship_flag	199998	2	N	199967
debt_settlement_flag	199998	2	N	194158

16 For the sake of the exercise, if the title appears 200 times, which is 1/1000 of the dataset, create a dummy variable for that title

```
[30]: len(df)/1000
```

```
[30]: 199.998
```

```
[31]: counts = pd.value_counts(df['emp_title'])
mask = df['emp_title'].isin(counts[counts > 200].index)
dummies = pd.get_dummies(df['emp_title'][mask], prefix='emp_title')
df = pd.concat([df, dummies], axis=1).fillna(0)
dummies
```

```
[31]:
```

	emp_title_-1	emp_title_account	executive	emp_title_account manager	\
1	0		0		0
2	0		0		0
8	0		0		0
10	0		0		0
14	0		0		0
...	
199986	0		0		0
199990	0		0		0
199993	0		0		0
199994	0		0		0
199995	0		0		0

	emp_title_accountant	emp_title_accounting manager	\
1	0		0
2	0		0
8	0		0
10	0		0
14	0		0
...

199986	0	0
199990	0	0
199993	0	0
199994	0	0
199995	0	0

	emp_title_administrative assistant	emp_title_administrator	\
1	0	0	
2	0	0	
8	0	0	
10	0	0	
14	0	0	
...	
199986	0	0	
199990	0	0	
199993	0	0	
199994	0	0	
199995	0	0	

	emp_title_analyst	emp_title_assistant manager	emp_title_associate	\
1	0	0	0	
2	0	0	0	
8	0	0	0	
10	0	0	0	
14	0	0	0	
...	
199986	0	0	0	
199990	0	0	0	
199993	0	0	0	
199994	0	0	0	
199995	0	0	0	

	...	emp_title_supervisor	emp_title_teacher	emp_title_team leader	\
1	...	0	0	0	
2	...	0	0	0	
8	...	0	0	0	
10	...	0	0	0	
14	...	0	0	0	
...	
199986	...	1	0	0	
199990	...	0	0	0	
199993	...	0	1	0	
199994	...	0	0	0	
199995	...	0	0	0	

	emp_title_tech	emp_title_technician	emp_title_truck driver	\
1	0	0	0	

2	0	0	1
8	0	0	0
10	0	0	0
14	0	0	0
...
199986	0	0	0
199990	0	0	0
199993	0	0	0
199994	0	0	1
199995	0	0	0

	emp_title_underwriter	emp_title_vice president	emp_title_vp \
1	0	0	0
2	0	0	0
8	0	0	0
10	0	0	0
14	0	0	0
...
199986	0	0	0
199990	0	0	0
199993	0	0	0
199994	0	0	0
199995	0	0	0

	emp_title_welder
1	0
2	0
8	0
10	0
14	0
...	...
199986	0
199990	0
199993	0
199994	0
199995	0

[72084 rows x 97 columns]

```
[32]: # list(df.select_dtypes(exclude='O').columns.values)
colnames = ['term',
'int_rate',
'grade',
'emp_length',
'dti',
'delinq_2yrs',
'inq_last_6mths',
```

```


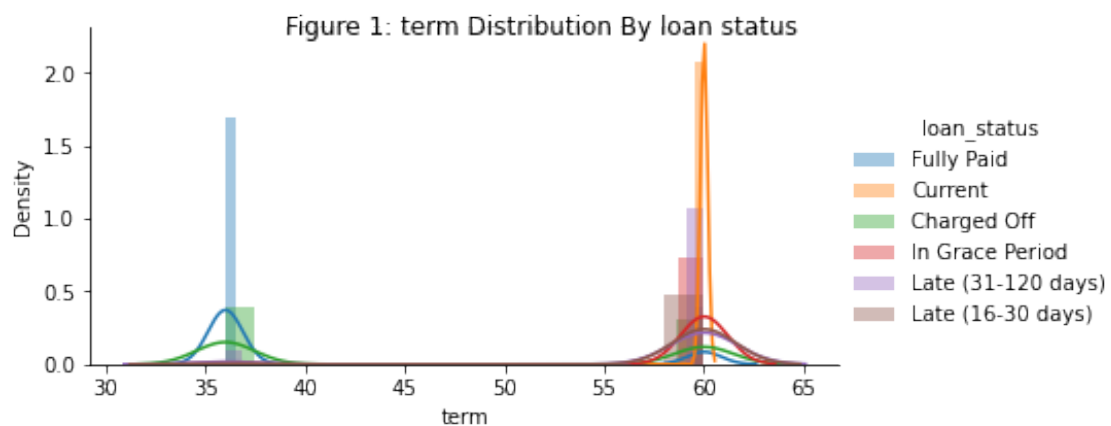
'mths_since_last_delinq',
'open_acc',
'pub_rec',
'total_rec_late_fee',
'last_pymnt_d',
'last_credit_pull_d',
'collections_12_mths_ex_med',
'mths_since_last_major_derog',
'policy_code',
'acc_now_delinq',
'acc_open_past_24mths',
'bc_open_to_buy',
'bc_util',
'chargeoff_within_12_mths',
'mo_sin_rcnt_rev_tl_op',
'mo_sin_rcnt_tl',
'mort_acc',
'mths_since_recent_bc',
'mths_since_recent_bc_dlq',
'mths_since_recent_inq',
'mths_since_recent_revol_delinq',
'num_accts_ever_120_pd',
'num_actv_bc_tl',
'num_actv_rev_tl',
'num_bc_sats',
'num_bc_tl',
'num_il_tl',
'num_op_rev_tl',
'num_rev_accts',
'num_rev_tl_bal_gt_0',
'num_sats',
'num_tl_120dpd_2m',
'num_tl_30dpd',
'num_tl_90g_dpd_24m',
'num_tl_op_past_12m',
'pct_tl_nvr_dlq',
'percent_bc_gt_75',
'pub_rec_bankruptcies',
'tax_liens']
len(colnames)

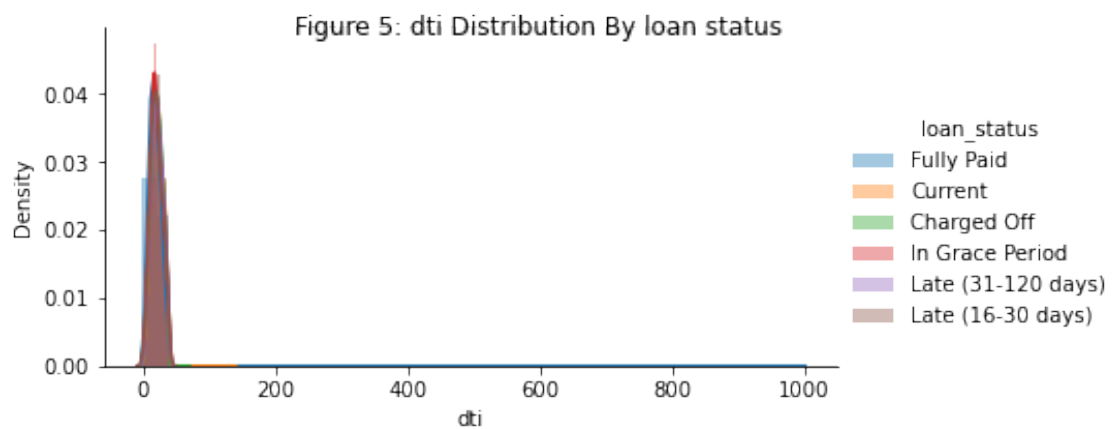

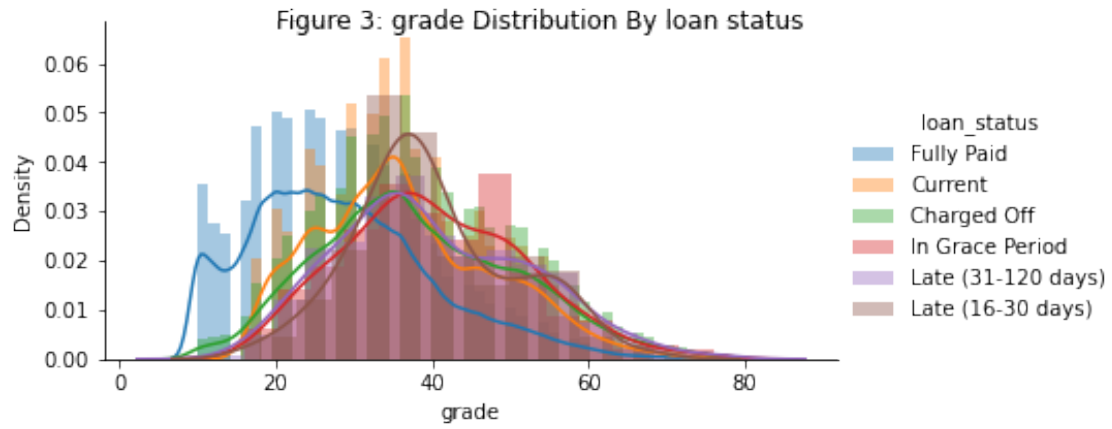
```

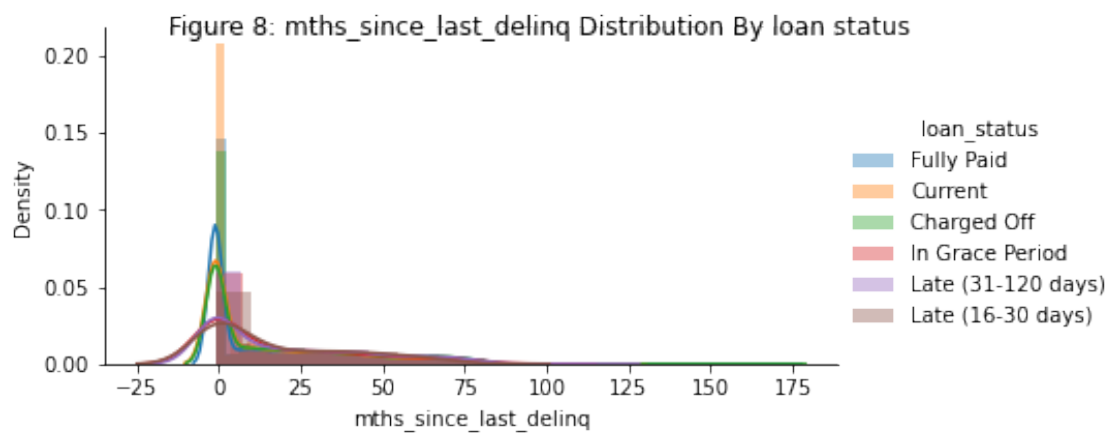
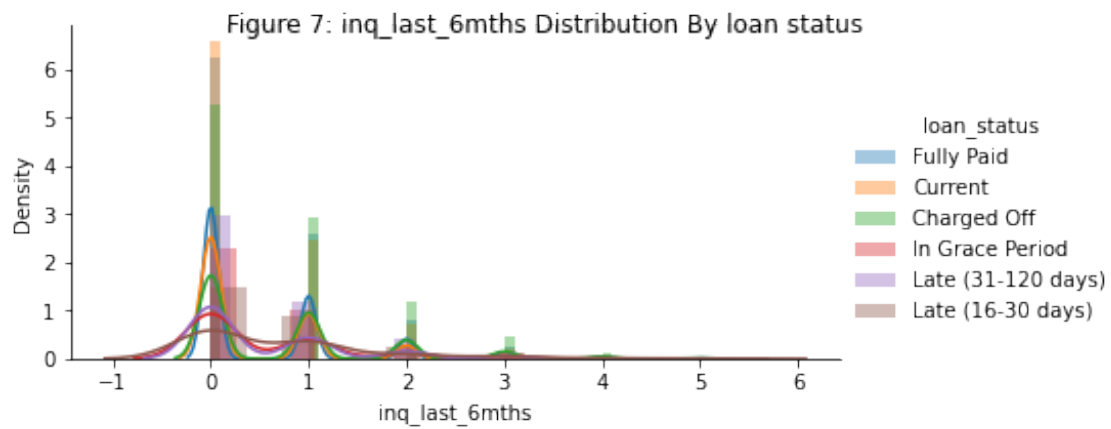
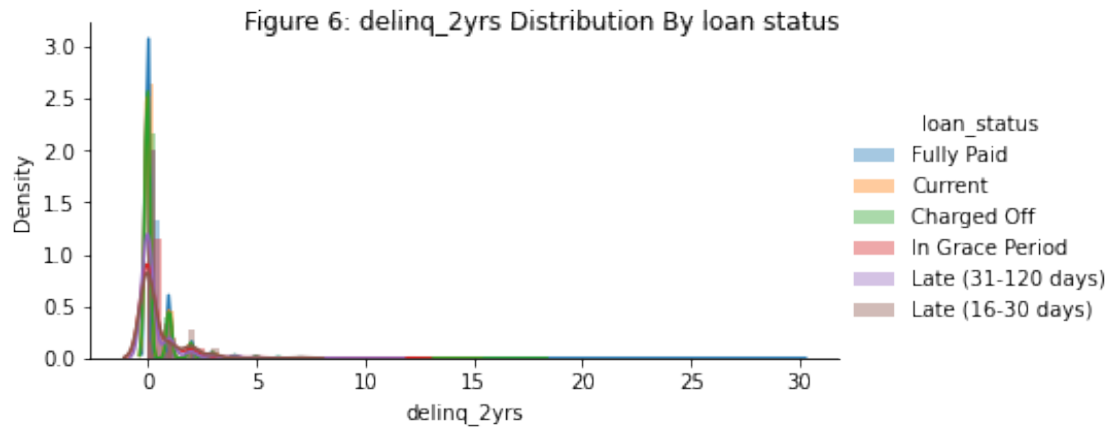
[32]: 46

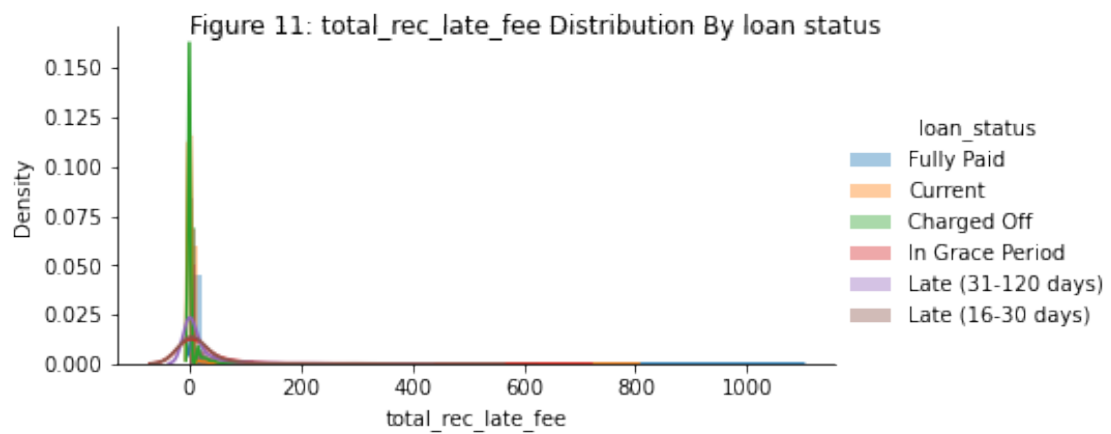
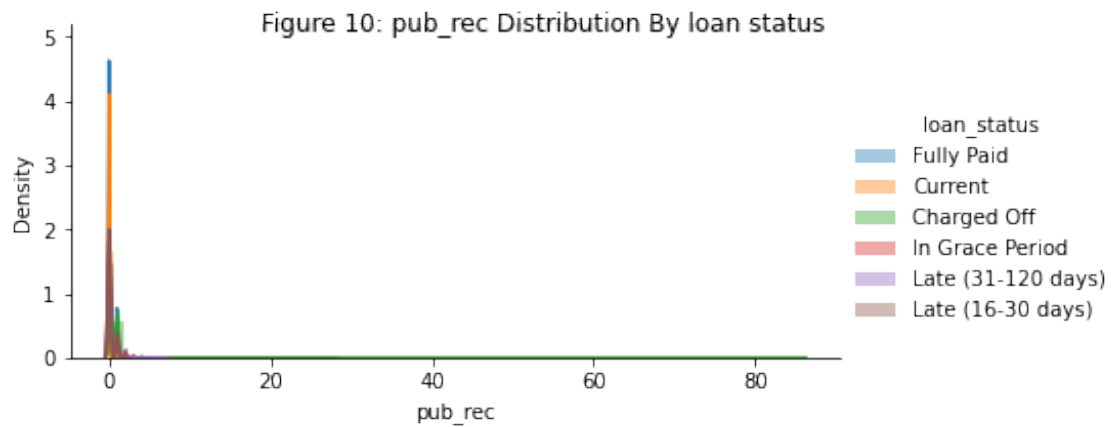
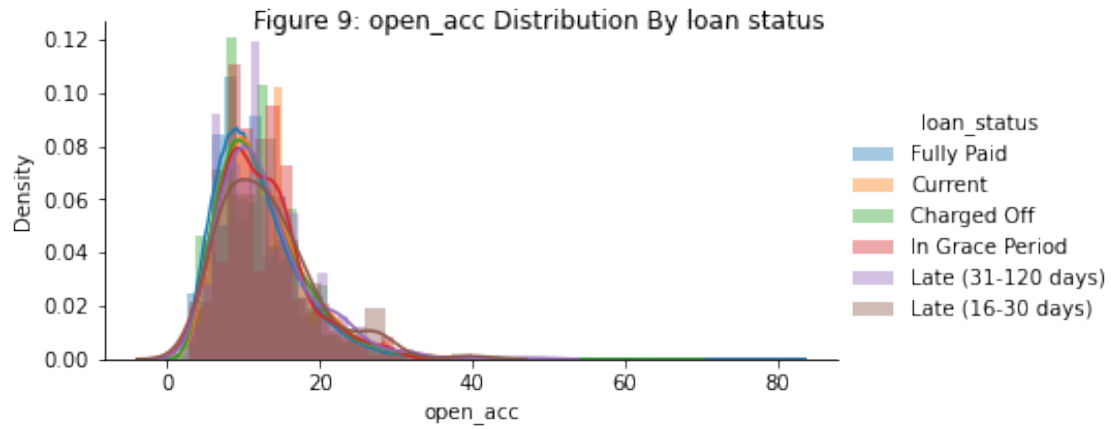
17 For the continuous variables, plot the distribution by loan status

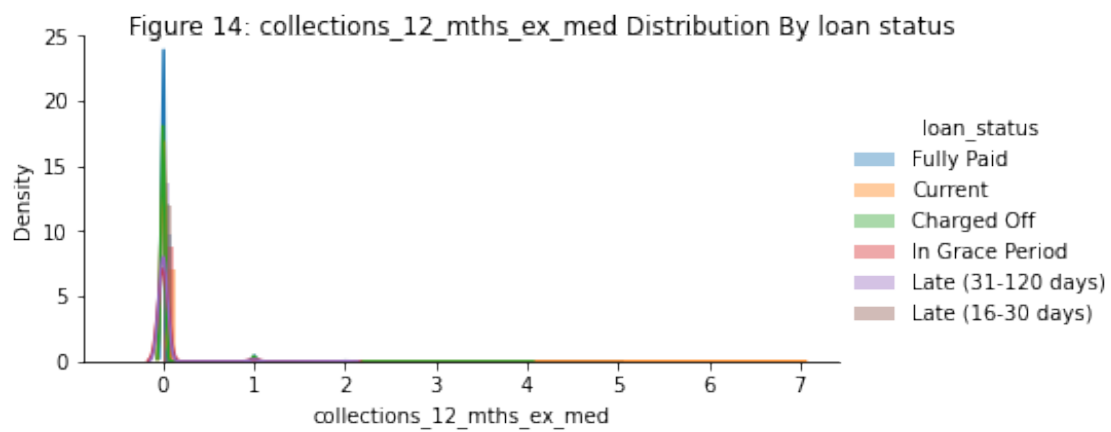

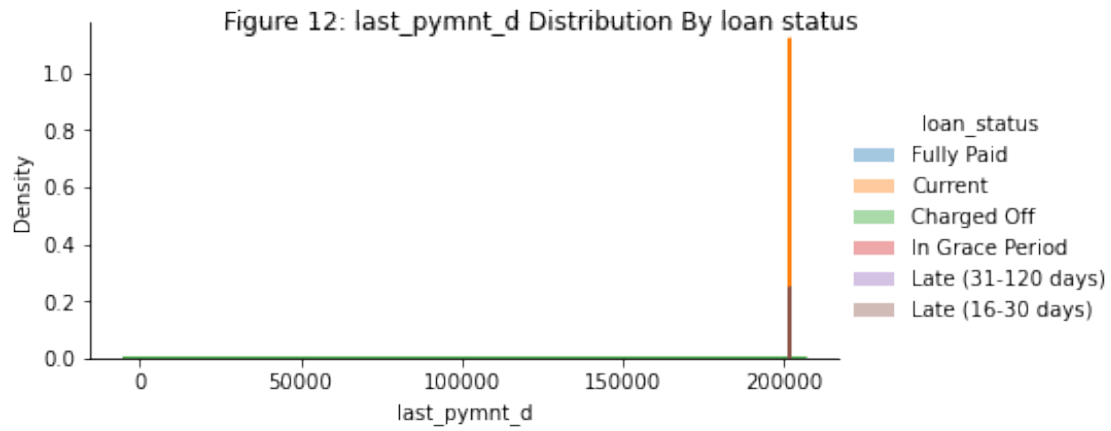
```
[33]: plt.rcParams['figure.max_open_warning']=40
# colnames=list(df.select_dtypes(exclude='O').columns.values)
for i in colnames[0:]:
    facet = sns.FacetGrid(df,hue='loan_status',aspect=2)
    facet.map(sns.distplot,i)
    facet.add_legend()
    facet.fig.suptitle(''.join(map(str, list(["Figure ",colnames.index(i)+1,"":_
↵",i," Distribution By loan status"]))))
plt.show()
```

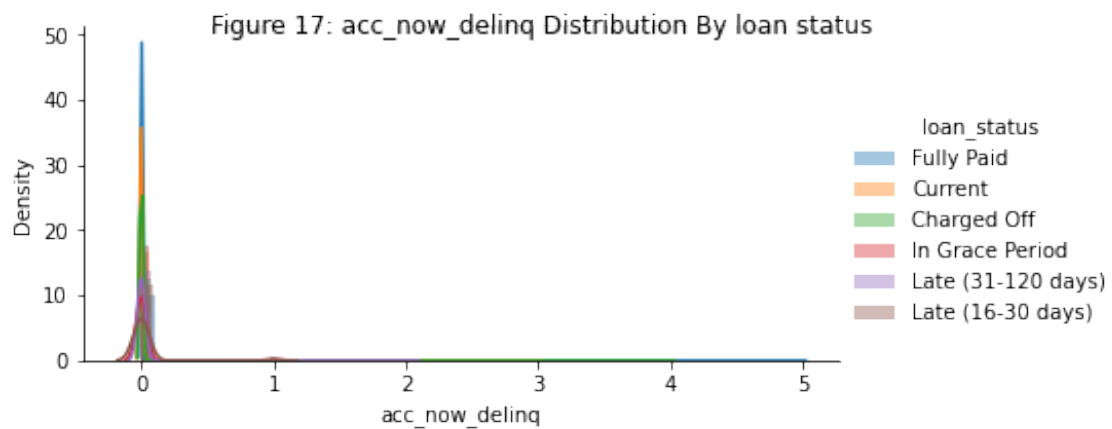
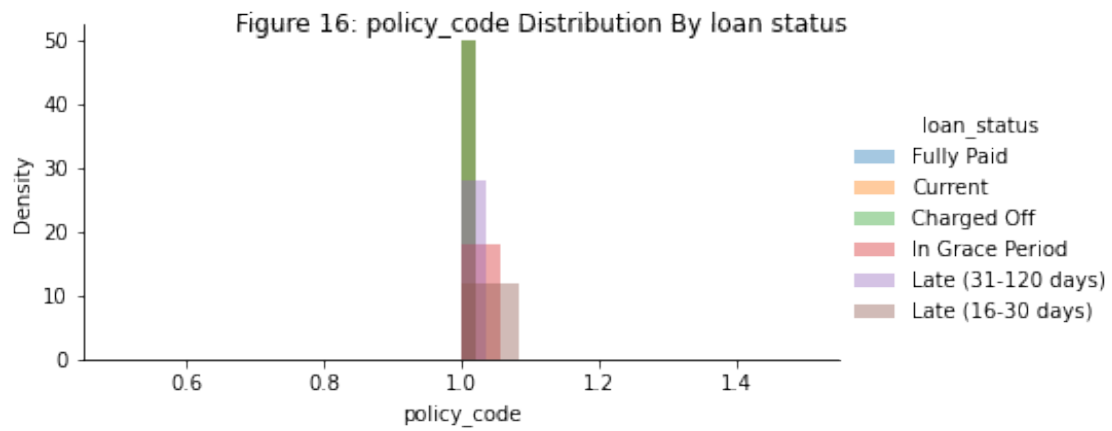
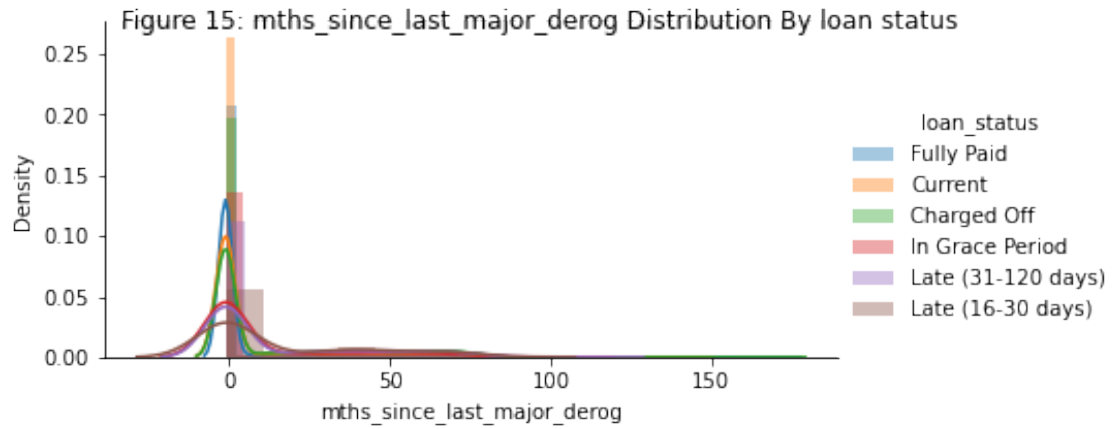



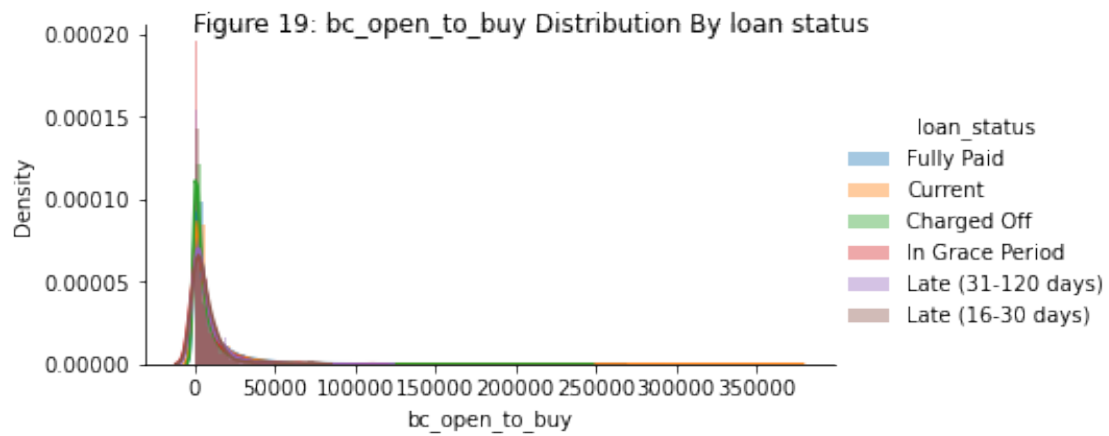
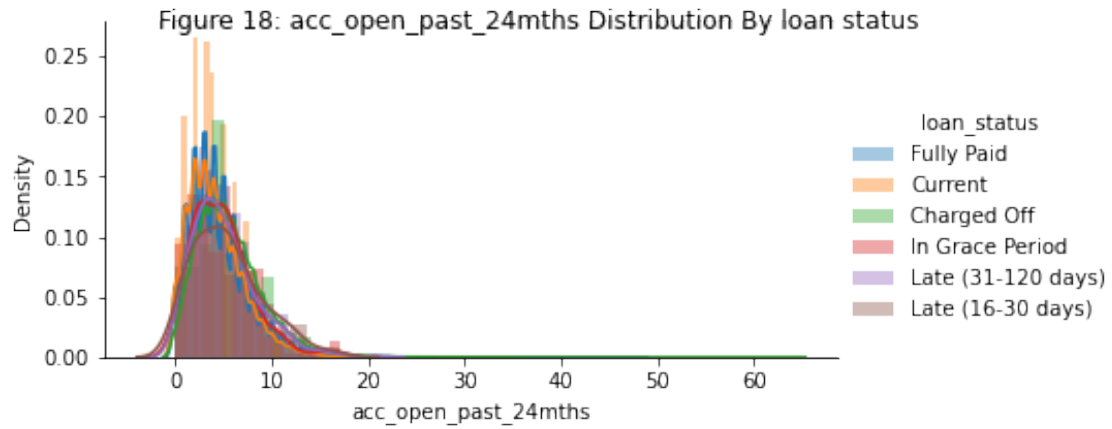


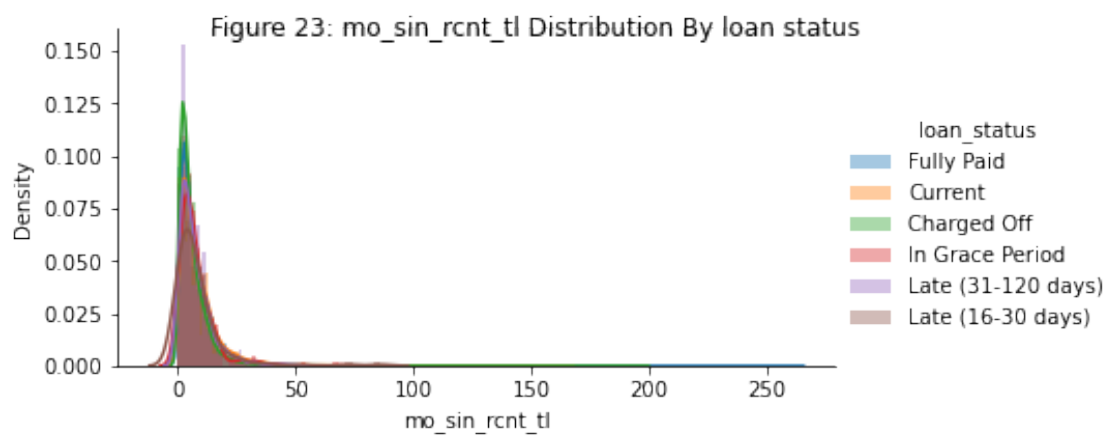
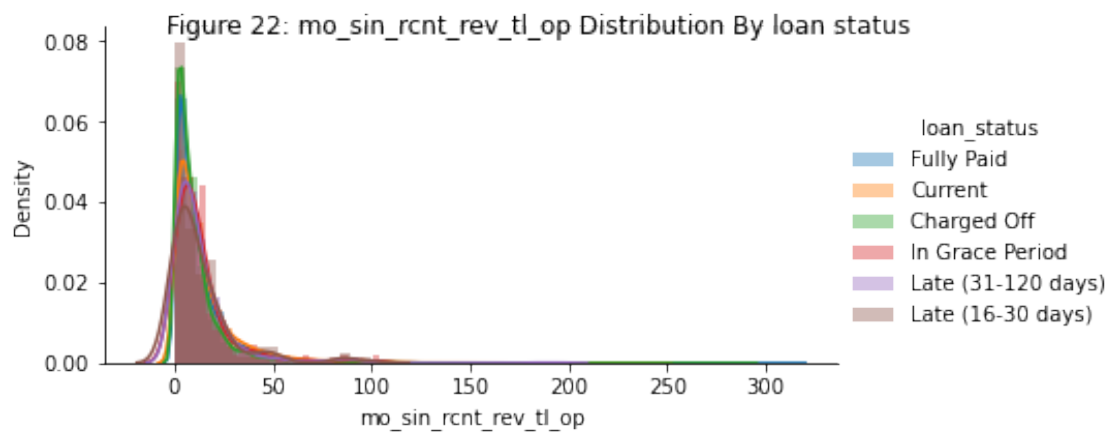
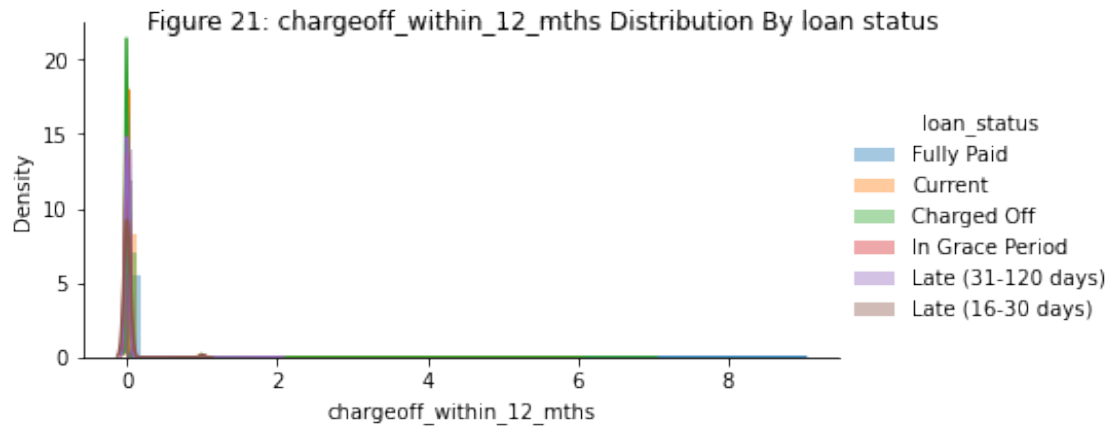


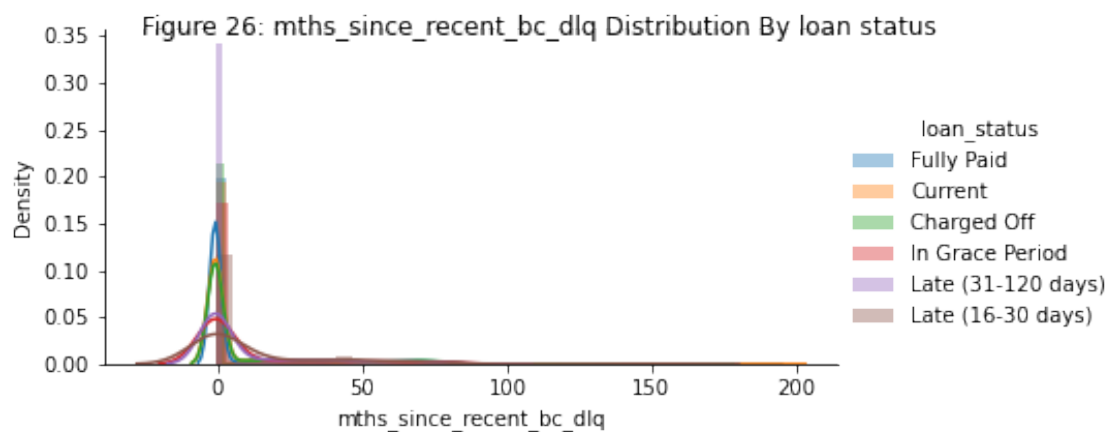
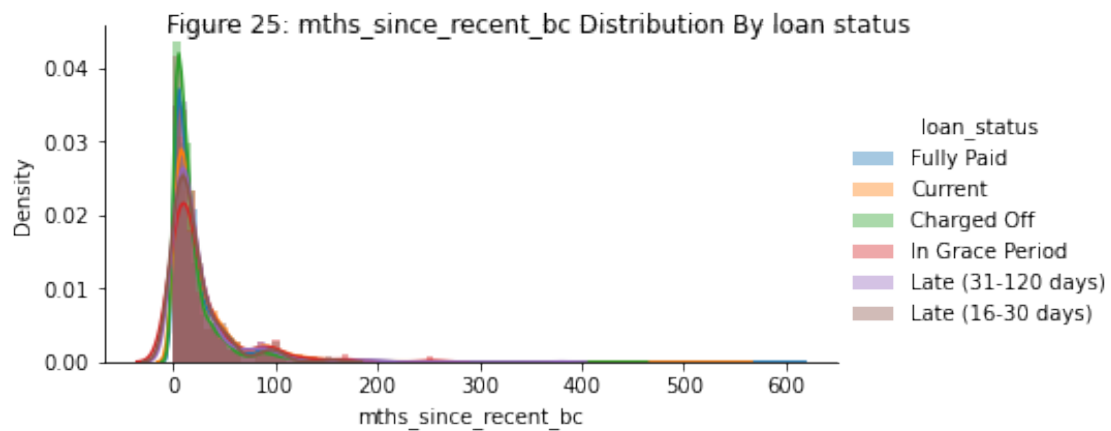
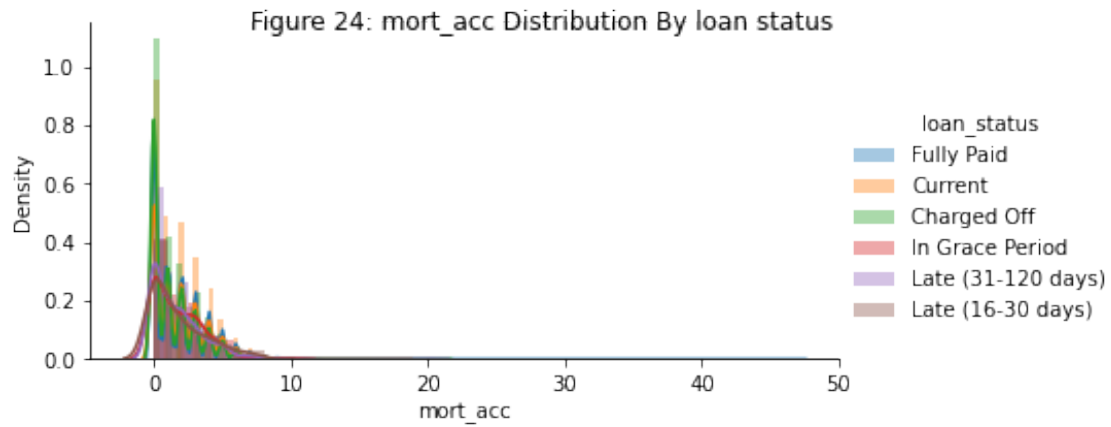


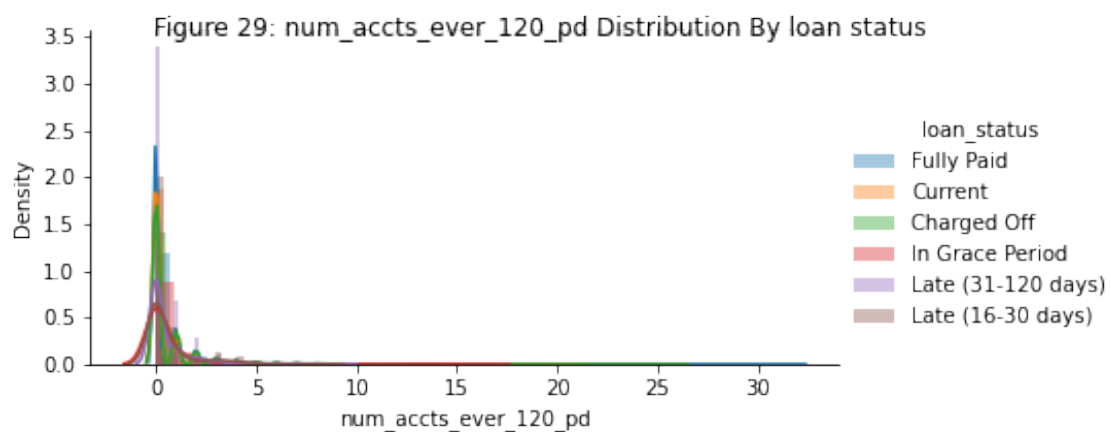
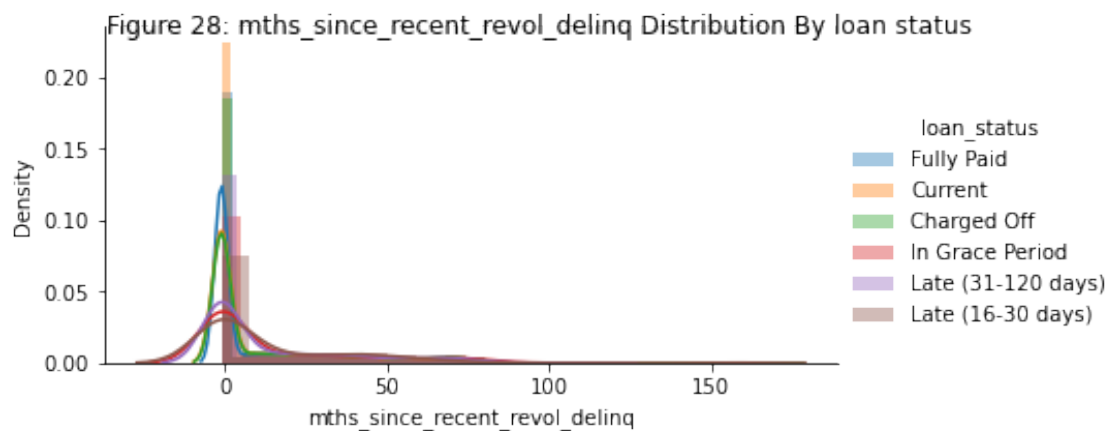
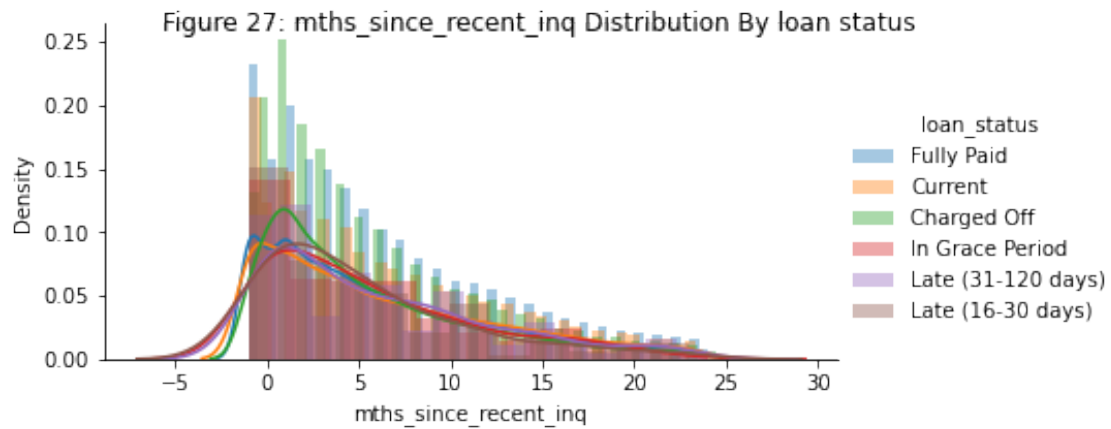


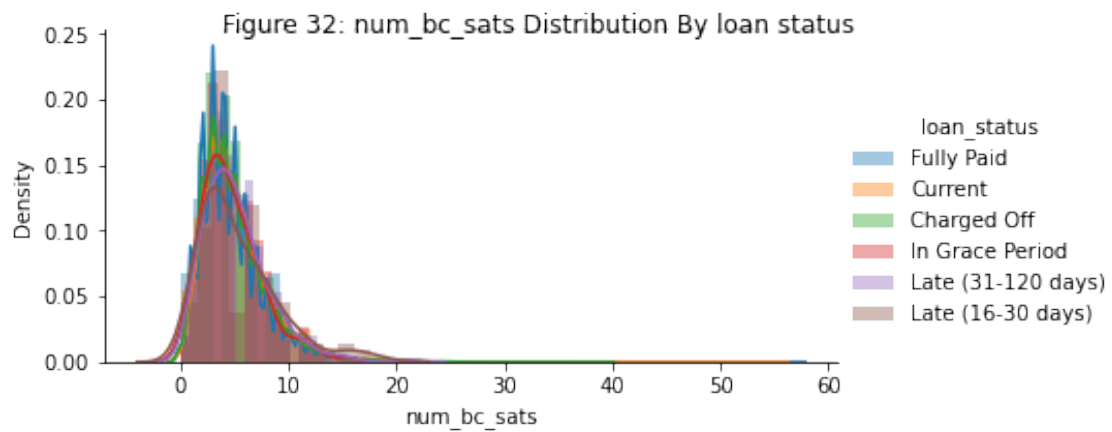
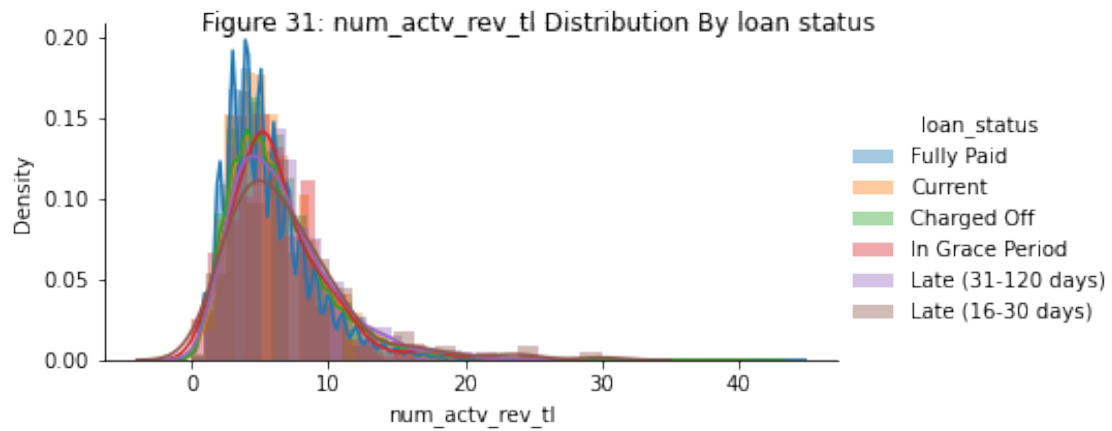
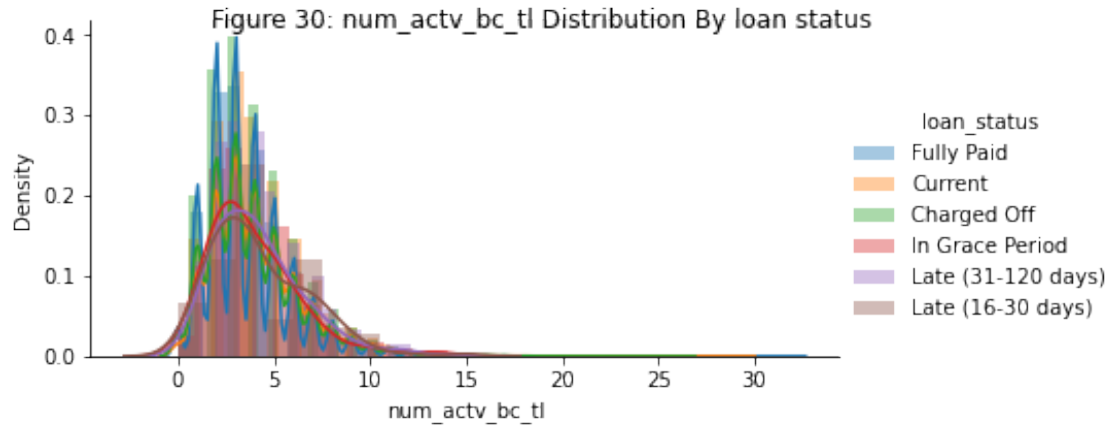




































```
[34]: # colnames=list(df.select_dtypes(include='O').columns.values)
colnames=['home_ownership',
'verification_status',
'loan_status',
'pymnt_plan',
'addr_state',
'initial_list_status',
'application_type',
'hardship_flag',
'debt_settlement_flag']
```

18 Visualization of categorical variables

```
[35]: plt.rcParams['figure.max_open_warning']=40
# colnames=list(df.select_dtypes(include='O').columns.values)
for i in colnames[0:]:
    ax = plt.axes()
    sns.countplot(x=i, data=df, ax = ax)
    ax.set_title(''.join(map(str, list(["Figure ",colnames.index(i)+47," ": ",i,"_
    ↪Distribution"]))))
    plt.show()
```




Figure 48: verification_status Distribution






Figure 49: loan_status Distribution





pymnt_plan	count
n	200000
y	10000

Bar chart showing the count of address states. The x-axis is labeled 'addr_state' and lists 26 states: PAD, UNF, DIR, NDIR, NOVA, Z, W, IE, SA, MA, CH, C, DM, A, H, VD, EN, D, AR, W, OK, MS, B, and BK. The y-axis is labeled 'count' and ranges from 0 to 25,000. The bars are colored in a gradient from red to purple. The highest count is for state 'NOVA' at approximately 27,000.

Figure 52: initial_list_status Distribution




Figure 53: application_type Distribution




Figure 54: hardship_flag Distribution






Figure 55: debt_settlement_flag Distribution




```
[36]: ax = plt.axes()
sns.countplot(x=df['emp_title'][mask], data=df, ax = ax)
ax.set_title(''.join(map(str, list(["Figure ", colnames.index(i)+63, ":",
↳ ", 'title', " Distribution"]))))
plt.show()
```



19 Separate the Target Variable from the rest of the dataset

```
[37]: y = df.pop('loan_status')
```

```
[38]: colnames.remove('loan_status')
```

20 Make dummy variables for select categorical variables

```
[39]: for i in colnames[0:]:
    # Fill missing data with the word "Missing"
    df[i].fillna("Missing", inplace=True)
    # Create array of dummies
    dummies = pd.get_dummies(df[i], prefix=i)
    # Update X to include dummies and drop the main variable
    df = pd.concat([df, dummies], axis=1)
    df.drop([i], axis=1, inplace=True)
```

```
df.drop('emp_title', axis=1, inplace=True)
```

- 21 Use minmaxscalar to scale all the dataset on a 0 to 1 scale for coefficient comparisons later for important features. Use stratified K fold with a set random state to shuffle the sampling set yet keep the results reproducible with the set random state

```
[40]: columns = df.columns
sc = MinMaxScaler(feature_range=(0, 1))
df = sc.fit_transform(df)
df = pd.DataFrame(df, columns=columns)

#x_train, x_test, y_train, y_test = train_test_split(df, y_int, test_size=0.33,
↳ random_state=2, stratify=y_int)
skf = StratifiedKFold(n_splits = 3, shuffle = True, random_state = 2)
```

- 22 Using a decision tree classified with default parameters, the metrics are already very good for the metrics of precision, recall, and f1-score. Recall is type 1 errors, which is when a person is predicted not to pay and pays. Precision is type 2 errors, which is when a person is predicted to pay and does not pay. F1 score is the balance between recall and precision, while accuracy is the overall reliability of the prediction

```
[41]: model2 = DecisionTreeClassifier()

for train_idx, test_idx in skf.split(df, y):
    y_score = model2.fit(df.loc[train_idx], y.loc[train_idx])
    print(classification_report(y.loc[test_idx], model2.predict(df.
↳ loc[test_idx])))
```

	precision	recall	f1-score	support
Charged Off	0.98	0.98	0.98	11697
Current	0.97	0.97	0.97	7546
Fully Paid	1.00	1.00	1.00	46997
In Grace Period	0.06	0.08	0.07	115
Late (16-30 days)	0.06	0.06	0.06	49
Late (31-120 days)	0.76	0.81	0.78	262
accuracy			0.99	66666
macro avg	0.64	0.65	0.64	66666
weighted avg	0.99	0.99	0.99	66666

	precision	recall	f1-score	support
Charged Off	0.98	0.98	0.98	11696
Current	0.97	0.97	0.97	7546
Fully Paid	0.99	0.99	0.99	46997
In Grace Period	0.10	0.10	0.10	116
Late (16-30 days)	0.02	0.02	0.02	49
Late (31-120 days)	0.75	0.76	0.75	262
accuracy			0.99	66666
macro avg	0.63	0.64	0.64	66666
weighted avg	0.99	0.99	0.99	66666

	precision	recall	f1-score	support
Charged Off	0.98	0.98	0.98	11697
Current	0.97	0.97	0.97	7545
Fully Paid	0.99	0.99	0.99	46997
In Grace Period	0.08	0.08	0.08	116
Late (16-30 days)	0.13	0.12	0.12	50
Late (31-120 days)	0.76	0.81	0.78	261
accuracy			0.99	66666
macro avg	0.65	0.66	0.66	66666
weighted avg	0.99	0.99	0.99	66666

23 The metrics for random forest classifier is also already very good

```
[42]: model3 = RandomForestClassifier(random_state=2, n_jobs=-1)

for train_idx, test_idx in skf.split(df, y):
    model3.fit(df.loc[train_idx], y.loc[train_idx])
    print(classification_report(y.loc[test_idx], model3.predict(df.
    ↪loc[test_idx])))
```

	precision	recall	f1-score	support
Charged Off	1.00	0.97	0.99	11697
Current	0.97	1.00	0.98	7546
Fully Paid	0.99	1.00	1.00	46997
In Grace Period	0.00	0.00	0.00	115
Late (16-30 days)	0.00	0.00	0.00	49
Late (31-120 days)	0.99	0.79	0.88	262

accuracy			0.99	66666
macro avg	0.66	0.63	0.64	66666
weighted avg	0.99	0.99	0.99	66666
	precision	recall	f1-score	support
Charged Off	1.00	0.97	0.98	11696
Current	0.97	0.99	0.98	7546
Fully Paid	0.99	1.00	1.00	46997
In Grace Period	0.00	0.00	0.00	116
Late (16-30 days)	0.00	0.00	0.00	49
Late (31-120 days)	0.99	0.71	0.83	262
accuracy			0.99	66666
macro avg	0.66	0.61	0.63	66666
weighted avg	0.99	0.99	0.99	66666
	precision	recall	f1-score	support
Charged Off	1.00	0.97	0.98	11697
Current	0.97	0.99	0.98	7545
Fully Paid	0.99	1.00	1.00	46997
In Grace Period	0.00	0.00	0.00	116
Late (16-30 days)	0.00	0.00	0.00	50
Late (31-120 days)	0.99	0.77	0.86	261
accuracy			0.99	66666
macro avg	0.66	0.62	0.64	66666
weighted avg	0.99	0.99	0.99	66666

24 Logistic regression is used to see if I could extrapolate if certain features could be removed from the model if they would cause overfitting because of too much noise from the dataset

```
[43]: model4 = LogisticRegression(penalty='l1', C=0.5, max_iter=100,
    ↪ solver='liblinear', multi_class='auto', random_state=2, n_jobs=-1)

for train_idx, test_idx in skf.split(df, y):
    model4.fit(df.loc[train_idx], y.loc[train_idx])
    print(classification_report(y.loc[test_idx], model4.predict(df.
    ↪ loc[test_idx])))
```

	precision	recall	f1-score	support
Charged Off	1.00	0.99	0.99	11697
Current	0.96	0.99	0.97	7546

Fully Paid	1.00	1.00	1.00	46997
In Grace Period	0.00	0.00	0.00	115
Late (16-30 days)	0.00	0.00	0.00	49
Late (31-120 days)	0.82	0.18	0.29	262
accuracy			0.99	66666
macro avg	0.63	0.53	0.54	66666
weighted avg	0.99	0.99	0.99	66666
	precision	recall	f1-score	support
Charged Off	1.00	0.99	0.99	11696
Current	0.95	0.99	0.97	7546
Fully Paid	0.99	1.00	1.00	46997
In Grace Period	1.00	0.01	0.02	116
Late (16-30 days)	0.00	0.00	0.00	49
Late (31-120 days)	0.87	0.15	0.25	262
accuracy			0.99	66666
macro avg	0.80	0.52	0.54	66666
weighted avg	0.99	0.99	0.99	66666
	precision	recall	f1-score	support
Charged Off	1.00	0.99	0.99	11697
Current	0.95	0.99	0.97	7545
Fully Paid	0.99	1.00	1.00	46997
In Grace Period	0.00	0.00	0.00	116
Late (16-30 days)	1.00	0.02	0.04	50
Late (31-120 days)	0.76	0.11	0.19	261
accuracy			0.99	66666
macro avg	0.79	0.52	0.53	66666
weighted avg	0.99	0.99	0.99	66666

25 here are the features that are certain, likely, possible, and unlikely towards contributing towards the prediction of the model

```
[44]: for b,c in enumerate(['certain','likely','possible','unlikely']):
      feature_importances = pd.Series(model4.coef_[b], index=df.columns)
      feature_importances.sort_values()[-30:].plot(kind="barh", figsize=(10,10),
          title = "Figure 64: Order of "+c+"
      ↪Features for Logistic Regression Classifier");
      plt.show()
```

Figure 64: Order of certain Features for Logistic Regression Classifier




Figure 64: Order of likely Features for Logistic Regression Classifier






Figure 64: Order of possible Features for Logistic Regression Classifier






26 For the important features that contribute to the model, only the ones above 0.0015 is included in the subset of the dataset

```
[45]: feature_importances = pd.Series(model3.feature_importances_, index=df.columns)
feature_importances.sort_values()[-30:].plot(kind="barh", figsize=(10,10),
                                             title = "Figure 65: Order of Features_
↳for Random Forest Classifier");
```

Figure 65: Order of Features for Random Forest Classifier



```
[46]: feature_importances.sort_values()[-29:]
```

```
[46]: annual_inc          0.001511
      total_rev_hi_lim    0.001511
      id                 0.001520
      revol_bal          0.001851
      dti                0.001858
      diff payment credit line 0.001879
      total_rec_late_fee    0.003359
      grade              0.006336
      int_rate           0.007127
      debt_settlement_flag_Y 0.009116
      debt_settlement_flag_N 0.011632
      installment        0.012405
      loan_amnt          0.012987
      funded_amnt_inv     0.015543
      total_pymnt_inv     0.021356
```

total_pymnt	0.023656
total_rec_int	0.026341
last_credit_pull_d	0.027415
term	0.027676
diff_credit_pull_payment	0.036426
total_rec_prncp	0.053586
last_fico_range_high	0.055023
last_pymnt_amnt	0.065148
last_fico_range_low	0.072411
collection_recovery_fee	0.083234
recoveries	0.086810
last_pymnt_d	0.089739
out_prncp	0.091004
out_prncp_inv	0.100441
dtype:	float64

27 Once the data is subsetting with just the important variables,
the model is ran again to see if the accuracy would increase
towards the important features

```
[47]: df = df[feature_importances.sort_values()[-29:].index]
df
```

[47]:	annual_inc	total_rev_hi_lim	id	revol_bal	dti	\
0	0.006111	0.005666	0.996940	0.001695	0.00691	
1	0.007222	0.068117	0.996179	0.013165	0.01706	
2	0.007000	0.008530	0.995985	0.004825	0.01178	
3	0.012222	0.041004	0.966360	0.004784	0.01806	
4	0.011604	0.020715	0.997954	0.013447	0.02637	
...	
199993	0.008778	0.015293	0.816843	0.014690	0.01674	
199994	0.009444	0.016024	0.817145	0.013206	0.01082	
199995	0.008889	0.014683	0.814654	0.012660	0.01462	
199996	0.006111	0.011272	0.792847	0.003521	0.01723	
199997	0.004444	0.021751	0.817567	0.018202	0.02629	
	diff payment	credit line	total_rec_late_fee	grade	int_rate	\
0		0.664560	8.649259e-12	0.382353	0.366286	
1		0.675050	8.649259e-12	0.294118	0.281791	
2		0.678264	8.649259e-12	0.235294	0.230672	
3		0.616036	8.649259e-12	0.411765	0.402619	
4		0.690402	8.649259e-12	0.735294	0.723701	
...		
199993		0.718687	8.649259e-12	0.294118	0.294466	
199994		0.634628	8.649259e-12	0.323529	0.311365	
199995		0.678264	3.188390e-02	0.294118	0.294466	

199996	0.648359	8.649259e-12	0.411765	0.394170
199997	0.784155	8.649259e-12	0.235294	0.239544

	debt_settlement_flag_Y	...	diff credit pull payment	\
0	0.0	...	0.521955	
1	0.0	...	0.882398	
2	0.0	...	0.743032	
3	0.0	...	0.510118	
4	0.0	...	0.731577	
...	
199993	0.0	...	0.592974	
199994	0.0	...	0.511264	
199995	0.0	...	0.511264	
199996	0.0	...	0.488354	
199997	0.0	...	0.580374	

	total_rec_prncp	last_fico_range_high	last_pymnt_amnt	\
0	0.102857	0.663529	0.003379	
1	0.705714	0.822353	0.025514	
2	0.571429	0.828235	0.435544	
3	0.545781	0.798824	0.022858	
4	0.297143	0.828235	0.278981	
...	
199993	0.114286	0.804706	0.010941	
199994	0.342857	0.845882	0.210610	
199995	0.600000	0.734118	0.018569	
199996	0.785714	0.822353	0.434945	
199997	0.200000	0.822353	0.006309	

	last_fico_range_low	collection_recovery_fee	recoveries	\
0	0.662722	0.0	0.0	
1	0.822485	0.0	0.0	
2	0.828402	0.0	0.0	
3	0.798817	0.0	0.0	
4	0.828402	0.0	0.0	
...	
199993	0.804734	0.0	0.0	
199994	0.846154	0.0	0.0	
199995	0.733728	0.0	0.0	
199996	0.822485	0.0	0.0	
199997	0.822485	0.0	0.0	

	last_pymnt_d	out_prncp	out_prncp_inv
0	0.999917	0.000000	0.000000
1	0.998638	0.000000	0.000000
2	0.999133	0.000000	0.000000
3	0.999959	0.700577	0.700577

4	0.998679	0.000000	0.000000
...
199993	0.999629	0.000000	0.000000
199994	0.998432	0.000000	0.000000
199995	0.999711	0.000000	0.000000
199996	0.998514	0.000000	0.000000
199997	0.999711	0.000000	0.000000

[199998 rows x 29 columns]

28 The metrics of the prediction is very good if not better with just the important features for less overfitting

```
[48]: model5 = RandomForestClassifier(random_state=2, n_jobs=-1)

for train_idx, test_idx in skf.split(df, y):
    model5.fit(df.loc[train_idx], y.loc[train_idx])
    print(classification_report(y.loc[test_idx], model5.predict(df.
    ↪loc[test_idx])))
```

	precision	recall	f1-score	support
Charged Off	1.00	0.99	0.99	11697
Current	0.97	0.99	0.98	7546
Fully Paid	1.00	1.00	1.00	46997
In Grace Period	0.00	0.00	0.00	115
Late (16-30 days)	0.00	0.00	0.00	49
Late (31-120 days)	0.96	0.81	0.88	262

accuracy			0.99	66666
macro avg	0.66	0.63	0.64	66666
weighted avg	0.99	0.99	0.99	66666

	precision	recall	f1-score	support
Charged Off	1.00	0.99	0.99	11696
Current	0.97	0.99	0.98	7546
Fully Paid	1.00	1.00	1.00	46997
In Grace Period	0.00	0.00	0.00	116
Late (16-30 days)	0.00	0.00	0.00	49
Late (31-120 days)	0.95	0.74	0.83	262


accuracy			0.99	66666
macro avg	0.65	0.62	0.63	66666
weighted avg	0.99	0.99	0.99	66666

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

Charged Off	1.00	0.99	0.99	11697
Current	0.97	0.99	0.98	7545
Fully Paid	1.00	1.00	1.00	46997
In Grace Period	1.00	0.01	0.02	116
Late (16-30 days)	0.50	0.04	0.07	50
Late (31-120 days)	0.96	0.78	0.86	261
accuracy			0.99	66666
macro avg	0.90	0.64	0.65	66666
weighted avg	0.99	0.99	0.99	66666

29 These are the important features ordered by importance

```
[49]: feature_importances = pd.Series(model5.feature_importances_, index=df.columns)
feature_importances.sort_values().plot(kind="barh", figsize=(10,10),
                                         title = "Figure 66: Order of Features_
                                         ↳for Random Forest Classifier");
```



30 The overall problem with random forest coefficients is that they are all positive so both debt settlement yes and no contribute towards the accurate prediction without realizing which feature is negatively or positively contributing towards the prediction, but these metrics is what the company should pull in regards to if a person will pay a loan and the model would determine what category of loan status the person would be in

```
[50]: feature_importances.sort_values()
```

```
[50]: annual_inc          0.001619
      grade              0.001634
      dti                0.001741
      revol_bal          0.001794
      total_rev_hi_lim    0.001795
      diff payment credit line 0.001846
      id                 0.002417
      total_rec_late_fee   0.002418
      int_rate            0.005609
      debt_settlement_flag_N 0.007780
      debt_settlement_flag_Y 0.009154
      total_pymnt         0.011034
      last_credit_pull_d   0.013124
      total_pymnt_inv      0.013795
      installment         0.014062
      funded_amnt_inv      0.015527
      total_rec_int        0.016861
      loan_amnt           0.017084
      term                0.019709
      diff credit pull payment 0.031190
      total_rec_prncp      0.049074
      last_pymnt_amnt      0.054992
      last_fico_range_low   0.081046
      last_fico_range_high  0.082950
      last_pymnt_d         0.089769
      collection_recovery_fee 0.098753
      recoveries           0.102494
      out_prncp            0.115178
      out_prncp_inv        0.135549
      dtype: float64
```

31 Fit the truncated dataset with hyperparameter optimization using gridsearch on the random forest model. The accuracy and other metrics are already very good.

```
[51]: hyperparameters = {'max_features': ['auto', 'sqrt', 'log2'],
                        'min_samples_split': [2,3,4]}
rfc = RandomForestClassifier(criterion='entropy', n_jobs=-1, random_state=2,
                             n_estimators=100)

search = GridSearchCV(rfc, hyperparameters, cv=3, scoring='accuracy',
                      verbose=3)

search.fit(df.loc[train_idx], y.loc[train_idx])

print(search.best_estimator_)
```

Fitting 3 folds for each of 9 candidates, totalling 27 fits

```
[CV 1/3] END max_features=auto, min_samples_split=2;, score=0.991 total time=
16.7s
[CV 2/3] END max_features=auto, min_samples_split=2;, score=0.993 total time=
11.1s
[CV 3/3] END max_features=auto, min_samples_split=2;, score=0.993 total time=
10.9s
[CV 1/3] END max_features=auto, min_samples_split=3;, score=0.992 total time=
10.8s
[CV 2/3] END max_features=auto, min_samples_split=3;, score=0.993 total time=
11.0s
[CV 3/3] END max_features=auto, min_samples_split=3;, score=0.992 total time=
10.9s
[CV 1/3] END max_features=auto, min_samples_split=4;, score=0.991 total time=
10.6s
[CV 2/3] END max_features=auto, min_samples_split=4;, score=0.993 total time=
10.8s
[CV 3/3] END max_features=auto, min_samples_split=4;, score=0.993 total time=
10.8s
[CV 1/3] END max_features=sqrt, min_samples_split=2;, score=0.991 total time=
10.5s
[CV 2/3] END max_features=sqrt, min_samples_split=2;, score=0.993 total time=
10.9s
[CV 3/3] END max_features=sqrt, min_samples_split=2;, score=0.993 total time=
10.7s
[CV 1/3] END max_features=sqrt, min_samples_split=3;, score=0.992 total time=
10.8s
[CV 2/3] END max_features=sqrt, min_samples_split=3;, score=0.993 total time=
10.7s
[CV 3/3] END max_features=sqrt, min_samples_split=3;, score=0.992 total time=
```



```

10.7s
[CV 1/3] END max_features=sqrt, min_samples_split=4;, score=0.991 total time=
10.1s
[CV 2/3] END max_features=sqrt, min_samples_split=4;, score=0.993 total time=
11.3s
[CV 3/3] END max_features=sqrt, min_samples_split=4;, score=0.993 total time=
11.3s
[CV 1/3] END max_features=log2, min_samples_split=2;, score=0.992 total time=
9.6s
[CV 2/3] END max_features=log2, min_samples_split=2;, score=0.993 total time=
9.4s
[CV 3/3] END max_features=log2, min_samples_split=2;, score=0.993 total time=
9.2s
[CV 1/3] END max_features=log2, min_samples_split=3;, score=0.992 total time=
9.5s
[CV 2/3] END max_features=log2, min_samples_split=3;, score=0.993 total time=
9.2s
[CV 3/3] END max_features=log2, min_samples_split=3;, score=0.992 total time=
9.5s
[CV 1/3] END max_features=log2, min_samples_split=4;, score=0.992 total time=
8.9s
[CV 2/3] END max_features=log2, min_samples_split=4;, score=0.993 total time=
9.2s
[CV 3/3] END max_features=log2, min_samples_split=4;, score=0.992 total time=
9.4s
RandomForestClassifier(criterion='entropy', min_samples_split=3, n_jobs=-1,
                        random_state=2)

```

32 The best parameter for random forest classifier is entropy, 3 sample split, and auto, which is then plugged in to retrieve the metrics with the subset important features of the dataset

```

[52]: rfc = search.best_estimator_

for train_idx, test_idx in skf.split(df, y):
    rfc.fit(df.loc[train_idx], y.loc[train_idx])
    print(classification_report(y.loc[test_idx], rfc.predict(df.loc[test_idx])))

```


	precision	recall	f1-score	support
Charged Off	1.00	0.99	0.99	11697
Current	0.97	0.99	0.98	7546
Fully Paid	1.00	1.00	1.00	46997
In Grace Period	0.00	0.00	0.00	115
Late (16-30 days)	0.00	0.00	0.00	49
Late (31-120 days)	0.97	0.79	0.87	262

accuracy			0.99	66666
macro avg	0.66	0.63	0.64	66666
weighted avg	0.99	0.99	0.99	66666
	precision	recall	f1-score	support
Charged Off	1.00	0.99	0.99	11696
Current	0.97	0.99	0.98	7546
Fully Paid	1.00	1.00	1.00	46997
In Grace Period	0.00	0.00	0.00	116
Late (16-30 days)	0.00	0.00	0.00	49
Late (31-120 days)	0.95	0.73	0.83	262
accuracy			0.99	66666
macro avg	0.65	0.62	0.63	66666
weighted avg	0.99	0.99	0.99	66666
	precision	recall	f1-score	support
Charged Off	1.00	0.99	0.99	11697
Current	0.97	0.99	0.98	7545
Fully Paid	1.00	1.00	1.00	46997
In Grace Period	0.00	0.00	0.00	116
Late (16-30 days)	0.50	0.02	0.04	50
Late (31-120 days)	0.96	0.76	0.85	261
accuracy			0.99	66666
macro avg	0.74	0.63	0.64	66666
weighted avg	0.99	0.99	0.99	66666

33 These are the most important features for the best model towards prediction of loan status

```
[53]: feature_importances = pd.Series(rfc.feature_importances_, index=df.columns)
feature_importances.sort_values().plot(kind="barh", figsize=(10,10),
title = "Figure 67: Order of Features_
↳for Random Forest Classifier");
```

Figure 67: Order of Features for Random Forest Classifier



- 34 Using a confusion matrix, here is the normalized and non-normalized labeling of the random forest classifier prediction label compared to the actual label. Type 1 error is a false positive depicted in the upper right of the diagonal while type 2 error is a false negative depicted in the bottom left of the diagonal. Given that loan companies make money off interest if people don't pay on time, they are more concerned if a person is labeled to not pay but pays on time, which would affect the bottom line and is a type 2 error. If a person is labeled as a someone who would pay but does not, that would be beneficial for the loan company to earn interest on the loan, which is a type 1 error.

```
[54]: def plot_confusion_matrix(cm, classes,
                                normalize=False,
                                title='Confusion matrix',
                                cmap=plt.cm.Blues):

    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
    plt.ylabel('True label')
    plt.xlabel('Predicted label')

# Compute confusion matrix
cnf_matrix = confusion_matrix(y.loc[test_idx], rfc.predict(df.loc[test_idx]))
```

```

np.set_printoptions(precision=2)


class_names = ['Charged Off', 'Current', 'Fully Paid', 'In Grace Period', 'Late (16-30 days)', 'Late (31-120 days)']

# Plot non-normalized confusion matrix
plt.figure(figsize = (15,5))
plt.subplot(1, 2, 1)
plot_confusion_matrix(cnf_matrix, classes=class_names, title='Figure 68A: Confusion matrix, without normalization')

# Plot normalized confusion matrix
plt.subplot(1, 2, 2)
plot_confusion_matrix(cnf_matrix, classes=class_names, normalize=True, title='Figure 68B: Normalized confusion matrix')

plt.tight_layout()
plt.show()

```



35 Even though the grace period and late are mislabeled by the confusion matrix, the target variable of Charged off is 0.99, along with fully paid and current at 0.99 and 1.00 in prediction accuracy, respectively

```

[55]: feature_list = list(df.columns)

# Import tools needed for visualization
from sklearn.tree import export_graphviz
import pydot

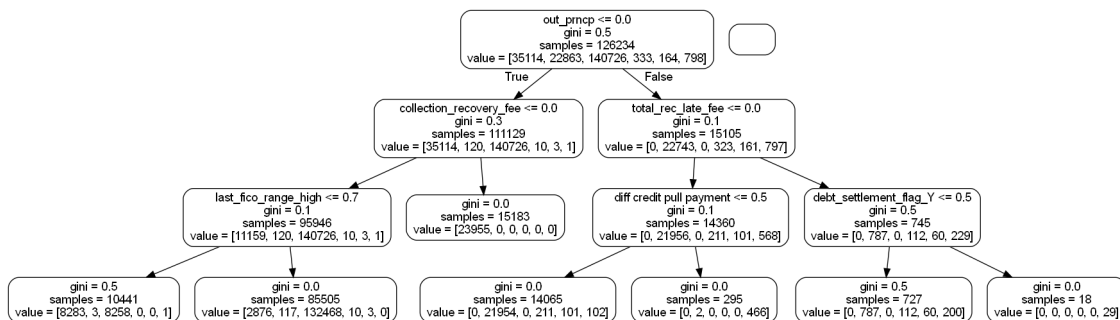
```

```
# Pull out one tree from the forest
tree = rfc.estimators_[5]
# Export the image to a dot file
export_graphviz(tree, out_file = 'tree.dot', feature_names = feature_list,
    rounded = True, precision = 1)
# Use dot file to create a graph
(graph, ) = pydot.graph_from_dot_file('tree.dot')
```

36 here is a graphic of random forest showing the decision trees in which a label is chosen from the dataset

```
[56]: # Limit depth of tree to 3 levels
rf_small = RandomForestClassifier(n_estimators=10, max_depth = 3)
rf_small.fit(df,y)
# Extract the small tree
tree_small = rf_small.estimators_[5]
# Save the tree as a png image
export_graphviz(tree_small, out_file = 'small_tree.dot', feature_names =
    feature_list, rounded = True, precision = 1)
(graph, ) = pydot.graph_from_dot_file('small_tree.dot')
graph.write_png('small_tree.png');
```

```
[57]: display(Image.open('small_tree.png'))
```



37 if a new people with the same data columns were to be added to the dataset with no label of their loan status, I could use the model to predict their label given that the accuracy of the model is already very high

```
[ ]:
```

```
[ ]:
```

[]: